

# CSE 673

# COMPUTATIONAL VISION

venu govindaraju  
deen dayal mohan


 University at Buffalo  
Department of Computer Science  
and Engineering  
School of Engineering and Applied Sciences

# Covid-19 Guidelines

- Effective Aug. 3, the University at Buffalo will require all students, employees and visitors – regardless of their vaccination status – to wear face coverings while inside campus buildings. This includes classrooms, hallways, libraries and other common spaces, as well as UB buses and shuttles.
- Students are expected to wear mask in class during lectures (unless you have a UB approved exception)
- Public Health Behavior Expectations <https://www.buffalo.edu/studentlife/who-we-are/departments/conduct/coronavirus-student-compliance-policy.html>


# Project Details

- Team of 3-4 depending upon final enrollment. 40% of the grade
- Set of topics provided by us
- Two presentations
- Reports of sample projects posted on Piazza
- Satisfies Masters Project requirement



# Agenda



- Computational Graph
- Backpropagation
- How is it implemented?
- Assignment 1



# COMPUTATIONAL GRAPH

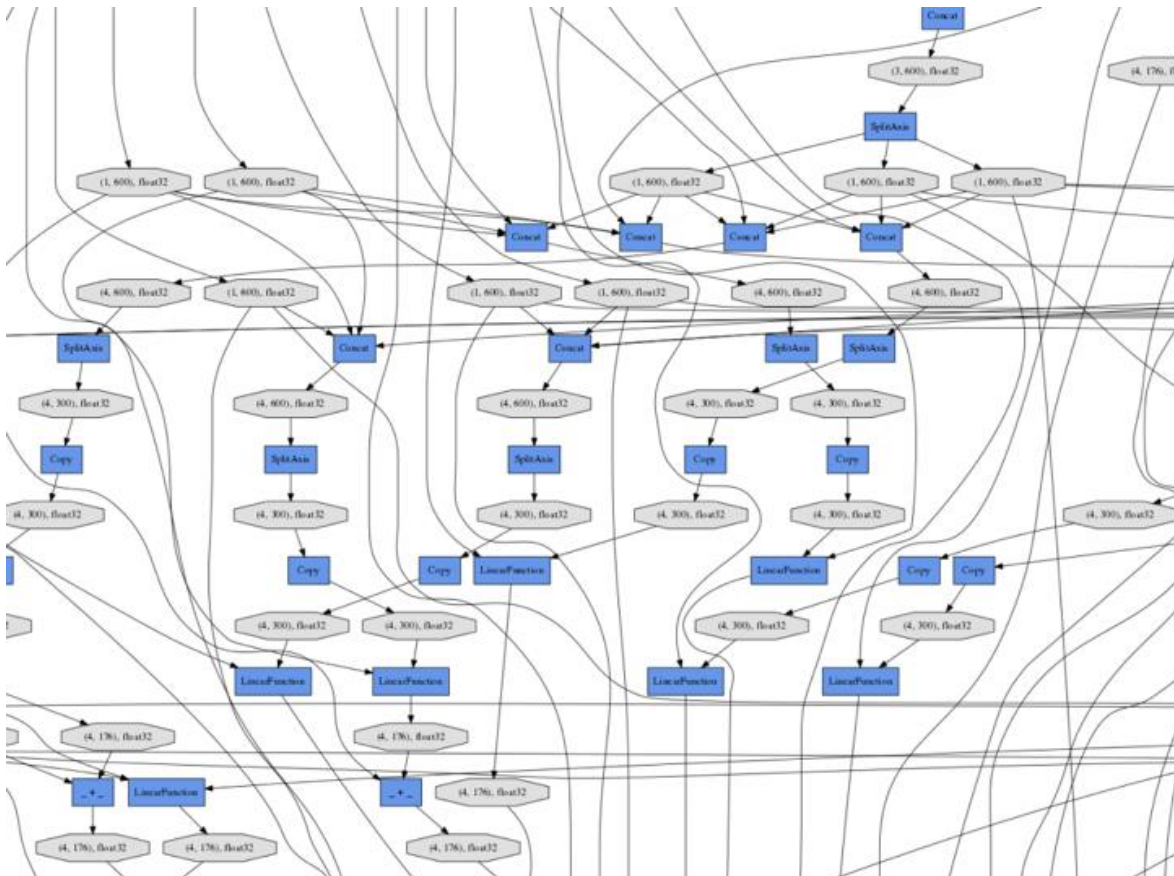
# Computational Graph

- Computational graphs are a nice way to think about mathematical expressions.
- For example consider  $O = (A + B) * C$ . The computational graph for the function would look like this:




# Computational Graph

- These computational graphs can get large.




Small section of the computational graph formed from an RNN based model

# Computational Graph



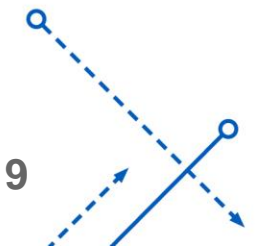

- These values are calculated when the computation is executed
- In the above diagram values flow from left to right







# Computational Graph

- The computational graph is directed and acyclic in nature
- The values are computed and fed in based on the topological order of the nodes
- The flow of values in the context of neural networks are called forward propagation




# Computational Graph

- Let us look at another example



# Computational Graph

- Now let us look at how to compute the influence of each input on the output





# Computational Graph

- What is the mathematical tool that we should use to find the influence of the inputs on the output ?
  - Partial derivatives to find the gradient


- So for the function  $O = (A + B) * C$  , we need to compute

- $$\frac{\partial O}{\partial A}, \frac{\partial O}{\partial B}, \frac{\partial O}{\partial C}$$



# Computational Graph


- These derivatives are not directly inferable. Let us see what are the derivatives that are directly inferable




- These derivatives are called local derivatives

# Computational Graph

- Let us think about the local derivatives a bit more




- F can be any mathematical function that is differentiable




# Computational Graph

- Let us look at an example




# Computational Graph

- Lets use real values instead




- The take away here should be, If input B is increased by x times it would change the output by a factor of 2x






# Computational Graph

- Let us look at another example



# Computational Graph


- Now we need to find the gradient with respect to the output. It happens that multiplying these gradients (applying chain rule) is the solution.



# Computational Graph


- Some additional details
- Imagine if the intermediate result of the function is used in two different computations FP1 and FP2
- In order to compute the gradient  $\frac{\partial \text{output}}{\partial A}$  and  $\frac{\partial \text{output}}{\partial B}$ , we will sum up all the gradients at F

- $$\frac{\partial \text{Output}}{\partial A} = \frac{\partial I}{\partial A} * \left( \frac{\partial P1}{\partial I} + \frac{\partial P2}{\partial I} \right)$$
$$\frac{\partial \text{Output}}{\partial B} = \frac{\partial I}{\partial B} * \left( \frac{\partial P1}{\partial I} + \frac{\partial P2}{\partial I} \right)$$



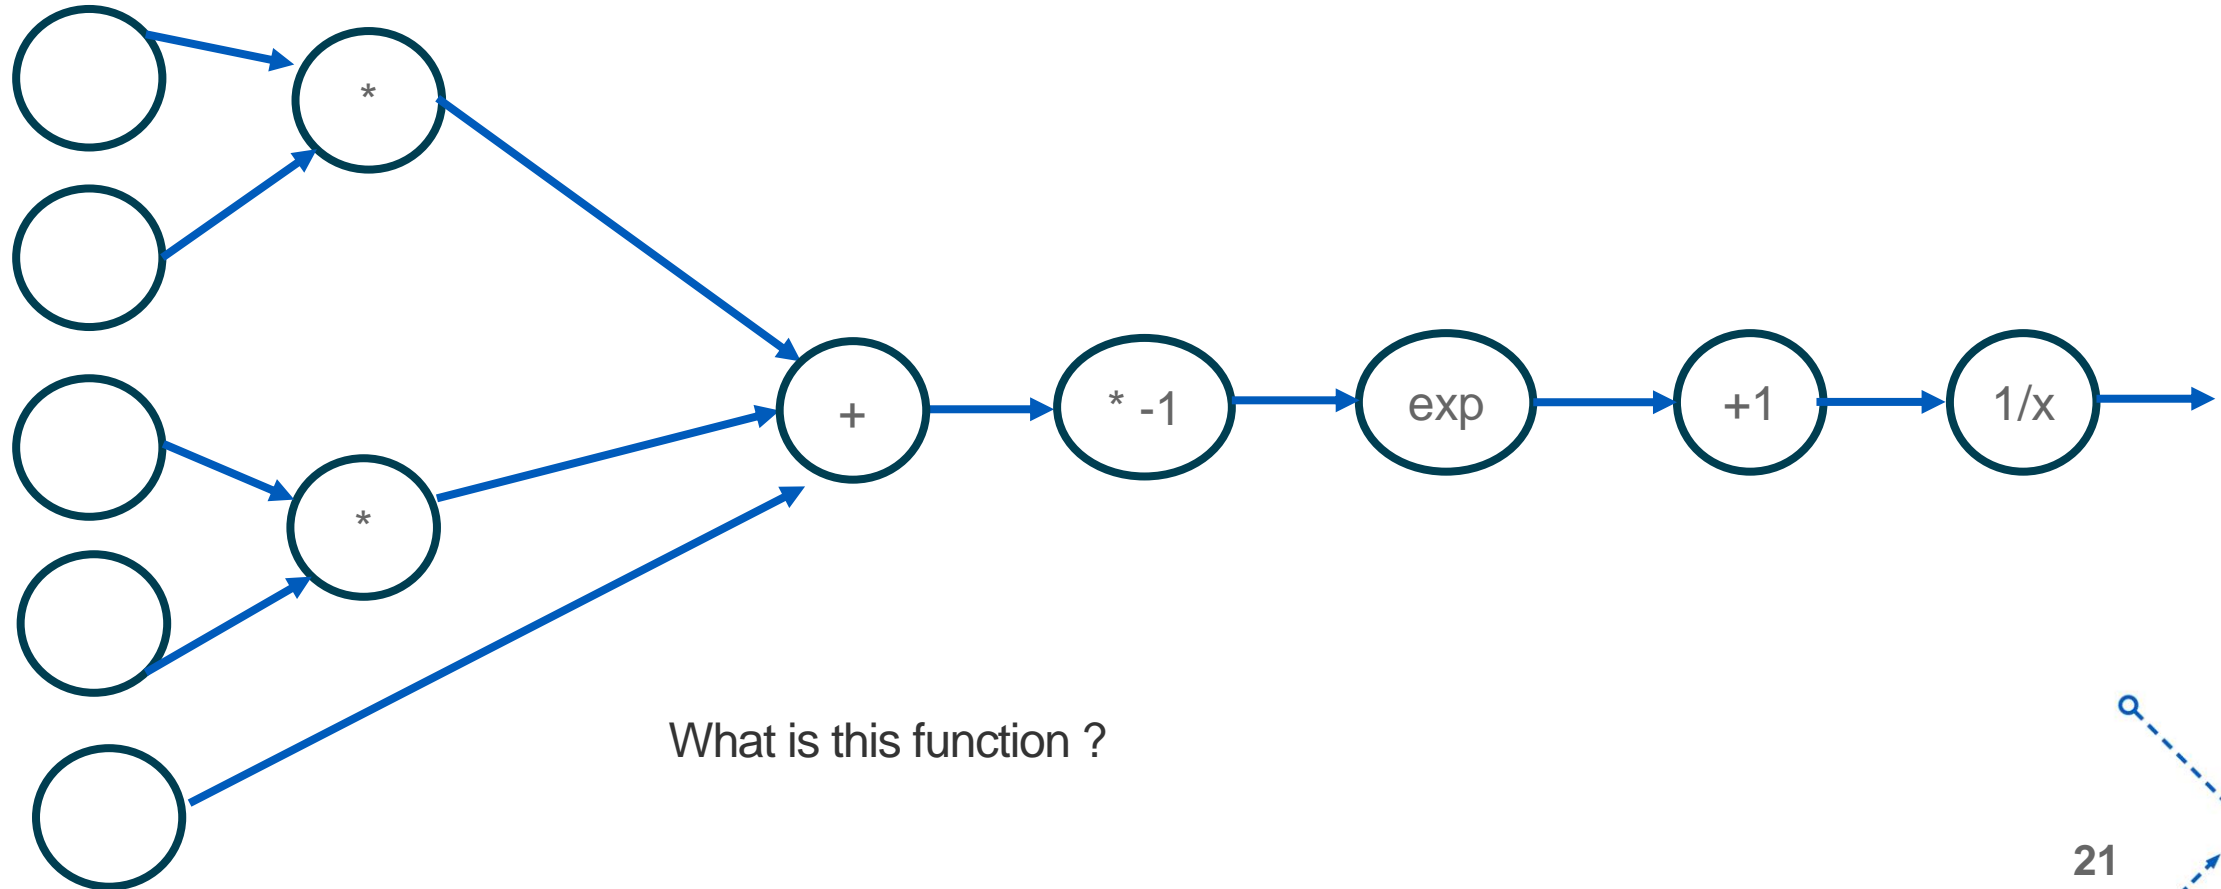
# Computational Graph

- Let us look at another example



# Computational Graph

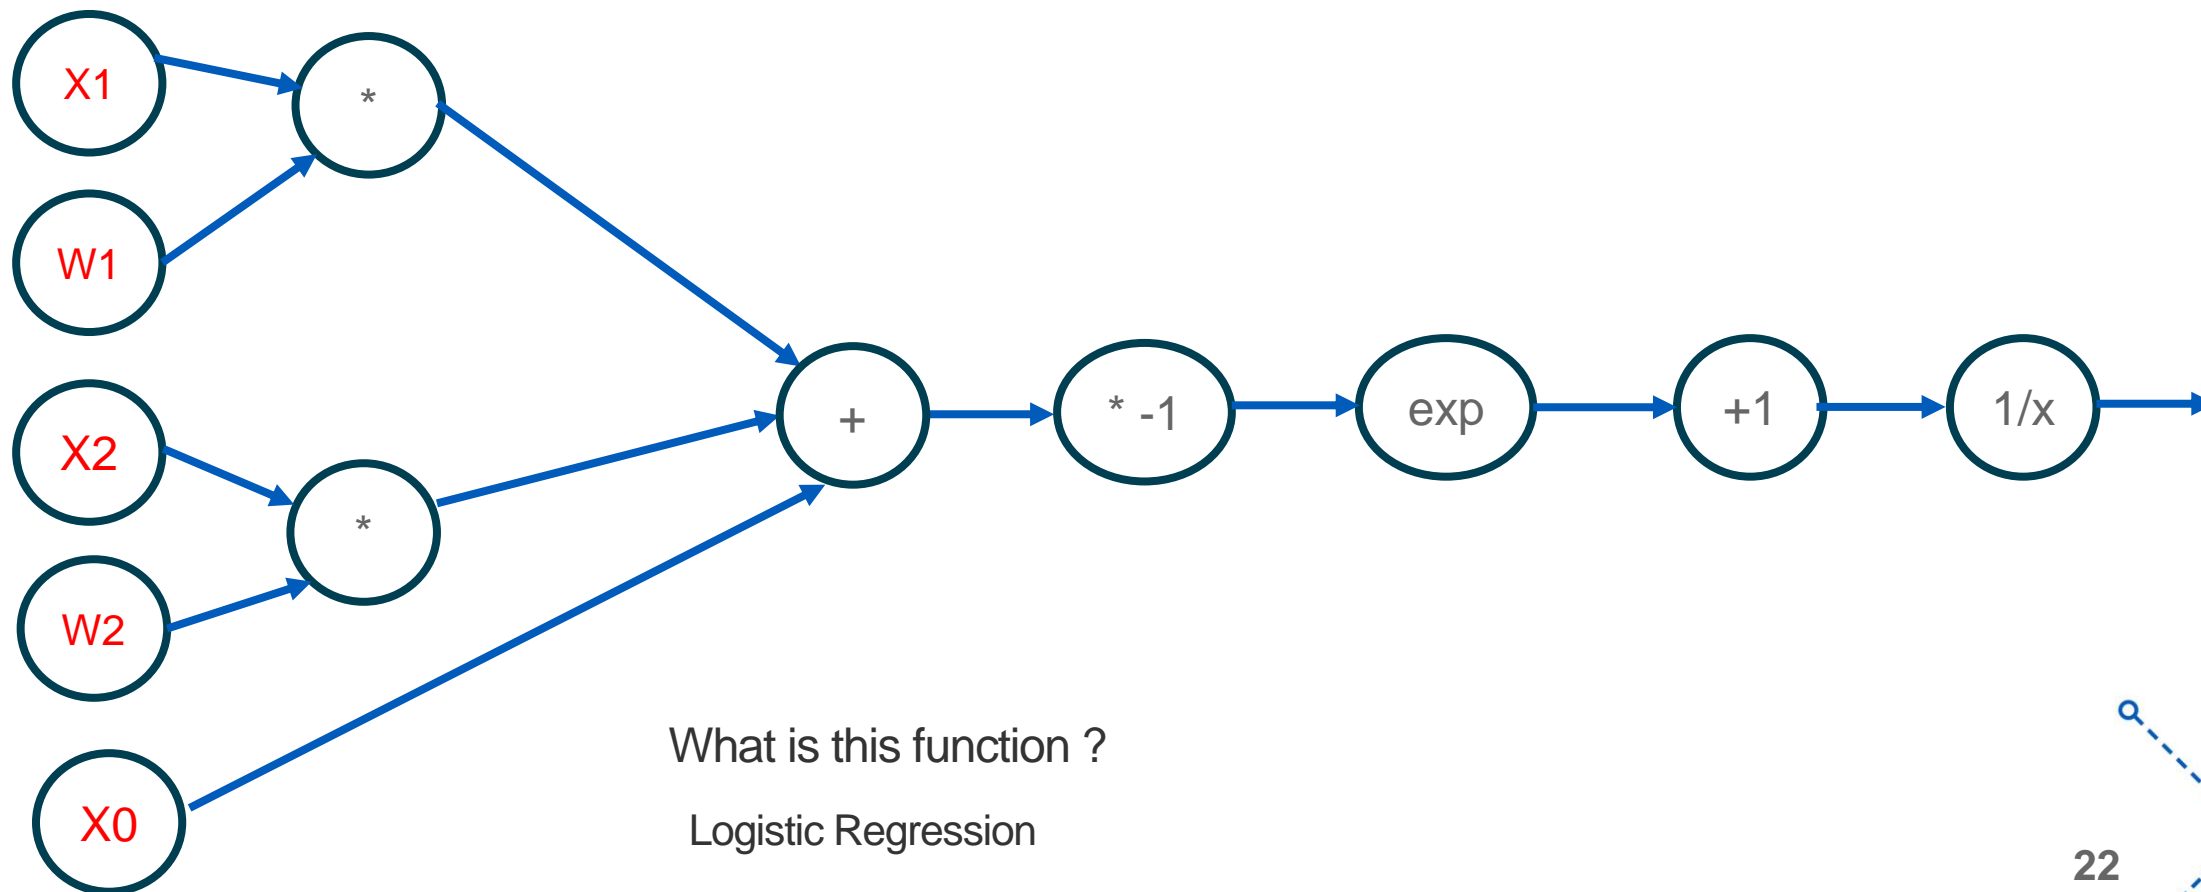
- Let us look at another example



What is this function ?

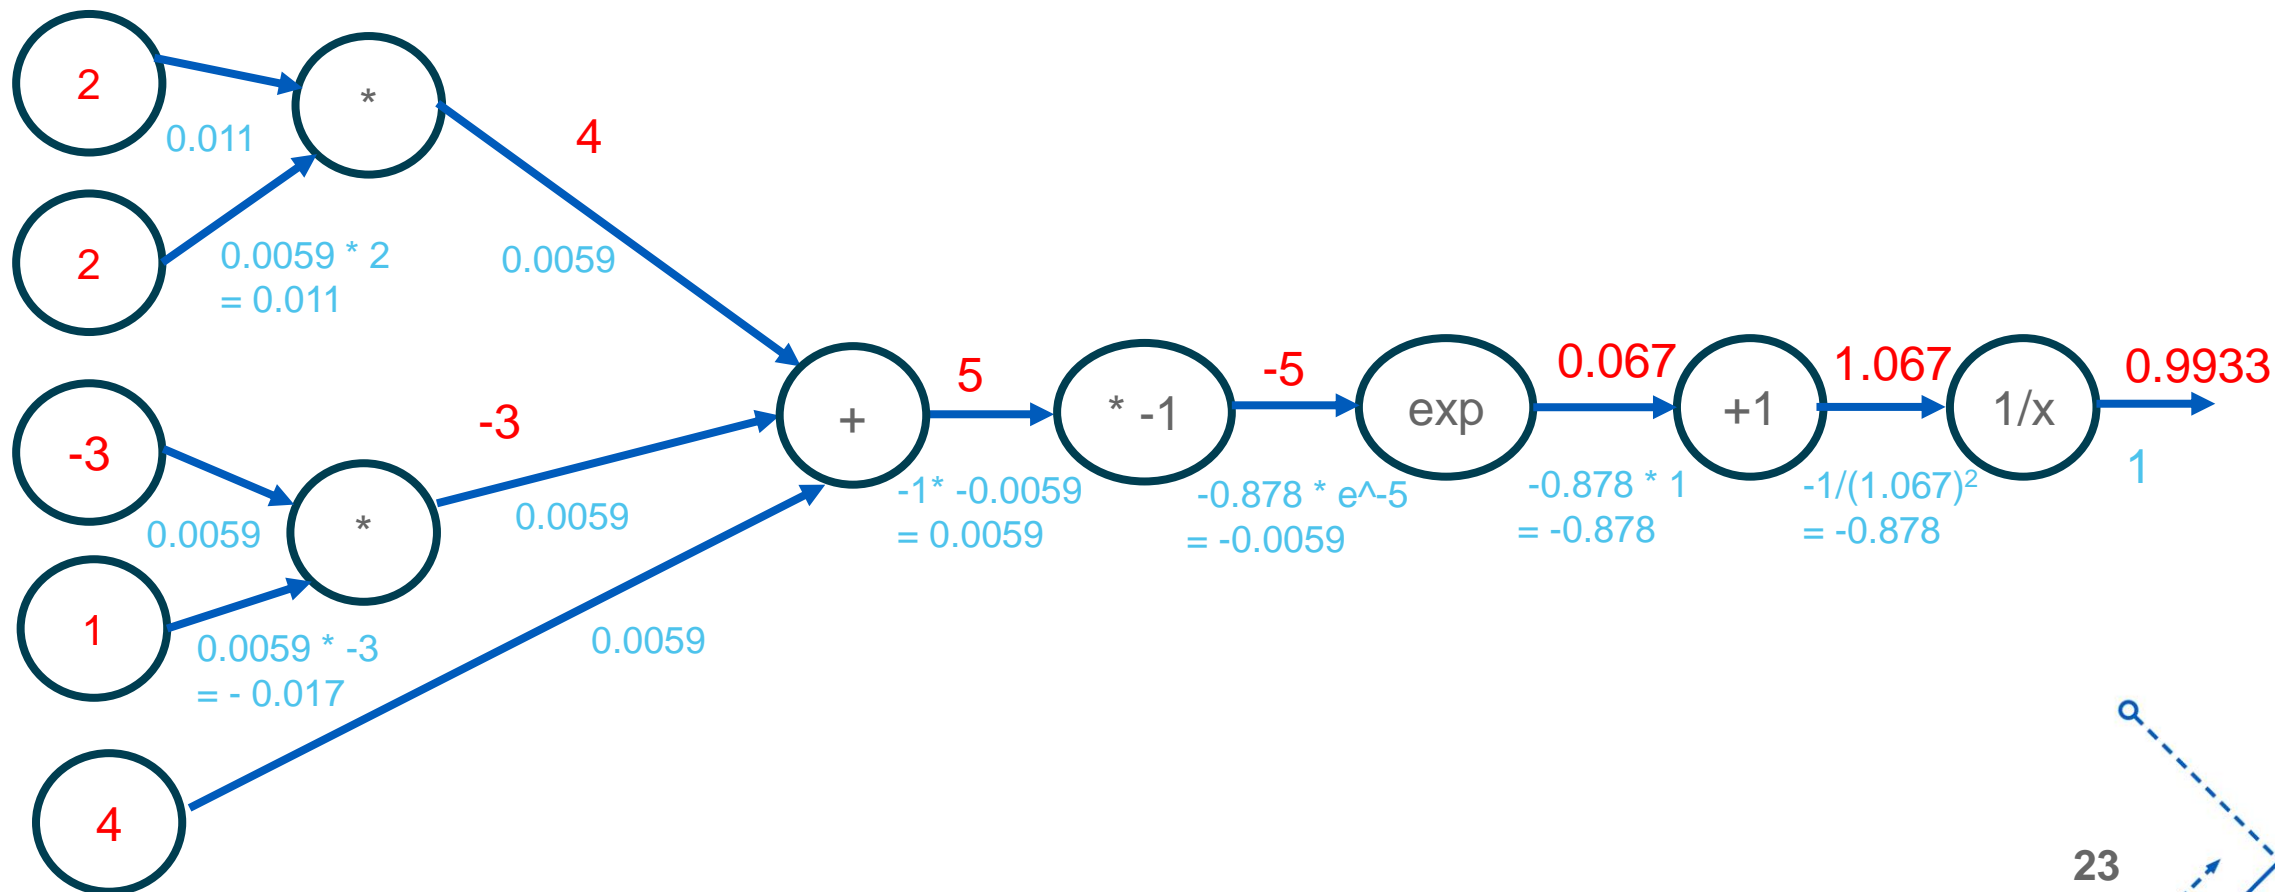
# Computational Graph

- Let us look at another example




# Computational Graph

- Let us look at another example




# Computational Graph






# Computational Graph



# Computational Graph


- The granularity of the nodes can be custom
- Need to find the derivative of the custom function
- For example, the derivative of  $\sigma(x)$  is

$$= \sigma(x)(1 - \sigma(x))$$




# Computational Graph

- This is often called a perceptron model
- When training the model, there would be an expected output  $Y$  associated with each data point
- Objective/Loss function is designed to estimate the error associated with the prediction and the expected value



# Computational Graph

- Some of the common loss functions you might know are MSE, Cross Entropy, Hinge etc.
- The derivative of the objective function with respect to the prediction is computed
- This is the first step in the backpropagation




# Computational Graph

- If we consider P as the prediction and MSE loss

- The loss formulation would be

$$J = \frac{1}{N} \sum_{i=1}^N (y - p)^2$$

- The first derivative found during the backward pass is  $\frac{\partial J}{\partial p}$



# Computational Graph

- Let us look at some code

```
class MaxNode:
```



```
    def forward(A,B):
```

```
        I = max(A,B)
        return I
```

```
    def backward(dI):
```

```
        dA = ..... # derivative of I with respect to A * dI
        dB = ..... # derivative of I with respect to B * dI
```

```
        return [dA,dB]
```



# Computational Graph


- We need one more thing to get it to work

```
class GraphNet:
    def __init__(self, graphnodes):
        self.graphnodes = graphnodes # Some representation of nodes in the computational graph
        ...
        ...
        ...

    def forward(input):
        for eachNode in topologically_sorted(self.graphnodes):
            eachNode.forward()
            ...
            ... # Logic to manipulate the output produced in the forward pass of each node
            ...
        return loss


    def backward():
        for eachNode in reversed(self.topologically_sorted(self.graphnodes)):
            eachNode.backward()
            ...
            ... # chain the gradients produced in the backward of each node
            ...

        return gradients
```



# Computational Graph


- The examples we saw until now consider  $X$  and  $W$  to be scalars
- But in real computations,  $X$  and  $W$  are vectors and the local gradients are Jacobians
- Jacobians are matrices
- The final computation for chained gradient will be a vector-matrix multiplication.






# Computational Graph

- Let us look at an example with matrices



# Artificial Neural Networks

- Inspired from the biological neural network in the brain
- Perceptron architecture was inspired from the working of a neuron




# Artificial Neural Networks

- Today there are many activations


## Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




## tanh

$$\tanh(x)$$




## ReLU

$$\max(0, x)$$



## Leaky ReLU

$$\max(0.1x, x)$$




## Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$


## ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



# Artificial Neural Networks

- Lets look at how the neural network is setup




- Why is this setup like this?



# Artificial Neural Networks

- A closer look



- Each column look like a stack of computational nodes. These stacks are called layers




# Artificial Neural Networks

- This neural network can be mathematically written as

$$\text{Output} = w_0^T \left( f_2 \left( w_2^T \left( f_1 (w_1^T x) \right) \right) \right)$$

- Where  $f_1, f_2$  are activation functions and  $w_1, w_2, w_0$  are the parameters of the network
- This can be modelled as a computational graph
- In order to improve computational efficiency, the basic computational unit can be the layer



# Artificial Neural Networks

- Let us look at some code

```
class GraphNet:

    def __init__(self, graphnodes):
        self.graphnodes = graphnodes # Some representation of layers in the computational graph
        ...
        ...
        ...

    def forward(input):
        for eachNode in topologically_sorted(self.graphnodes):
            eachNode.forward()
            ...
            ... # logic to manipulate the output produced in the forward pass of each layer
            ...
        return loss

    def backward():
        for eachNode in reversed(self.topologically_sorted(self.graphnodes)):
            eachNode.backward()
            ...
            ... # chain the gradients produced in the backward of each layer
            ...

        return gradients
```



# Artificial Neural Networks


- But now instead of the nodes we have layers

```
class BaseLayer:

    def __init__(self,):
        ...
        ... # initialize some parameters common to all the layers
        ...

    def forward():
        pass

    def backward():
        pass
```






# Artificial Neural Networks

- Then inherit the base class for implementing a layer with some specific functionality

```
class DenseLayer(BaseLayer):
    def __init__(self, input_param):
        ...
        ... # initialize the base class constructor
        ... # initialize layer specific parameters

    def forward(input):
        ...
        ... # perform forward propagation
        ...
        return layerOutput

    def backward(dz):
        ...
        ... # perform backprop
        ...
        return gradient
```




# Artificial Neural Networks

Consider the derivative of error function with respect to the weight from the hidden unit  $j$  to unit  $l$  where  $j = 1, 2, \dots, M + 1$  and  $l = 1, \dots, K$ :

$$\begin{aligned}\frac{\partial J_i}{\partial w_{lj}^{(2)}} &= \frac{\partial J_i}{\partial o_l} \frac{\partial o_l}{\partial b_l} \frac{\partial b_l}{\partial w_{lj}^{(2)}} \\ &= \delta_l z_j\end{aligned}$$


On the other hand, the derivative of error function with respect to the weight from the  $p^{th}$  input to hidden unit  $j$  where  $p = 1, 2, \dots, D + 1$  and  $j = 1, \dots, M$  can be computed as follow:

$$\begin{aligned}\frac{\partial J_i}{\partial w_{jp}^{(1)}} &= \sum_{l=1}^K \frac{\partial J_i}{\partial o_l} \frac{\partial o_l}{\partial b_l} \frac{\partial b_l}{\partial z_j} \frac{\partial z_j}{\partial a_j} \frac{\partial a_j}{\partial w_{jp}^{(1)}} \\ &= \sum_{l=1}^K \delta_l w_{lj}^{(2)} (1 - z_j) z_j x_p \\ &= (1 - z_j) z_j \left( \sum_{l=1}^k \delta_l w_{lj}^{(2)} \right) x_p\end{aligned}$$



# Artificial Neural Networks

- Once the gradients with respect to the necessary parameters are computed, update step is performed
- $$w_{new} = w_{old} - \eta \frac{\partial J}{\partial w_{old}}$$
- Here  $\eta$  is the learning rate
- All the parameters are updated after the gradients are computed
- We will look at other methods of parameter update in a later lecture




# Assignment 1

- Create a Toy Deep Learning framework called “PyTorchNano”.
- PyTorchNano should support the following layers
  - Dense Layer
  - Convolutional Layer
  - Max and Average Pooling Layer
  - Flatten Layer
- It should also support the following activations
  - ReLU
  - Softmax
  - Sigmoid




# Assignment 1

- The user should be able to train the model with following loss functions:
  - Cross-Entropy
  - Hinge Loss
  - MSE
- The framework should allow the user to create a custom network architecture (not more than 10 layers deep) and train it on the MNIST dataset.
- The Implementation of Layers and Activation functions should be using classes having the forward and backward method.
- The Layers should be customizable (number of hidden neurons, filter size etc.)
- The framework should also provide the capability to measure the accuracy of the model.



# Assignment 1

- Framework should provide capability to visualize filter responses of a convolutional layer
- Bonus points will be awarded if the framework has Batch Norm Layer and Dropout Layer implemented.
- The implementation should not use any Deep Learning frameworks like Pytorch, Keras , Tensorflow etc.
- More specific details will be mentioned in the assignment document, which will be released before next class
- The numbers of members in a team might be from 1-2 depending upon final enrollment



ANY  
QUESTIONS  
?

# References

- ❑ This lecture is inspired from cse 231n <https://www.youtube.com/watch?v=i94OvYb6noo&t=2051>
- ❑ <http://neuralnetworksanddeeplearning.com/chap5.html>
- ❑ <https://mlcourse-ub.readthedocs.io/en/latest/>
- ❑ <https://developer.nvidia.com/blog/recursive-neural-networks-pytorch/>
- ❑ <http://cs231n.stanford.edu/>
- ❑ [https://www.google.com/imgres?imgurl=https%3A%2F%2Fmiro.medium.com%2Fmax%2F1400%2F1\\*Di4V69e4gC16ooF6PZPt-A.png&imgrefurl=https%3A%2F%2Ftowardsdatascience.com%2Feverything-you-need-to-know-about-neural-networks-and-backpropagation-machine-learning-made-easy-e5285bc2be3a&tbnid=PFKzBNejYXM4hM&vet=12ahUKEwism4altO3yAhVrqnlEhSHUCNkQMyhDegQIARBF..i&docid=OXeL--Z4fRwo6M&w=1250&h=1057&q=neural%20networks%20with%20math&hl=en&client=firefox-b-1-d&ved=2ahUKEwism4altO3yAhVrqnlEhSHUCNkQMyhDegQIARBF](https://www.google.com/imgres?imgurl=https%3A%2F%2Fmiro.medium.com%2Fmax%2F1400%2F1*Di4V69e4gC16ooF6PZPt-A.png&imgrefurl=https%3A%2F%2Ftowardsdatascience.com%2Feverything-you-need-to-know-about-neural-networks-and-backpropagation-machine-learning-made-easy-e5285bc2be3a&tbnid=PFKzBNejYXM4hM&vet=12ahUKEwism4altO3yAhVrqnlEhSHUCNkQMyhDegQIARBF..i&docid=OXeL--Z4fRwo6M&w=1250&h=1057&q=neural%20networks%20with%20math&hl=en&client=firefox-b-1-d&ved=2ahUKEwism4altO3yAhVrqnlEhSHUCNkQMyhDegQIARBF)

